

▼ Credit Card Fraud Detection Project

Keeping customers money safe is primary job of any bank. Any fraud with credit card end up with loss to bank in terms of money, market reputation and customer trust.

Dataset: <https://www.kaggle.com/mlg-ulb/creditcardfraud>

The dataset contains transactions made by credit cards in September 2013 by European cardholders. This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

▼ 1. Importing the Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

▼ 2. Loading the dataset

```
from google.colab import drive
drive.mount('/content/drive/')
```

Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount("/content/d

```
df = pd.read_csv('/content/drive/MyDrive/projectData/creditcard.csv')
df.head()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	..
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	

5 rows × 31 columns



▼ 3. Data Preprocessing

```
df.shape
```

(284807, 31)

Here we have total 31 columns and 284807 rows of data.

```
df.describe()
```

	Time	V1	V2	V3	V4	V5	
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	9.604066e-16	1.487313e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1.332271e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-7.682956e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-2.741871e-01
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	3.985649e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7.330163e+01

8 rows × 31 columns



Handle missing values

```
# Cheking missing values in columns
df_missing_values = df.isnull().sum()
df_missing_values
```

Time	0
V1	0
V2	0
V3	0
V4	0
V5	0
V6	0
V7	0
V8	0
V9	0
V10	0
V11	0
V12	0
V13	0
V14	0
V15	0
V16	0
V17	0
V18	0
V19	0
V20	0
V21	0
V22	0
V23	0
V24	0
V25	0
V26	0
V27	0
V28	0
Amount	0


```
Class    0  
dtype: int64
```

Data distribution analysis

```
classes = df['Class'].value_counts()  
classes
```

```
0    284315  
1      492  
Name: Class, dtype: int64
```

```
sns.countplot(x='Class', data=df)  
plt.title('Number of fraud vs non-fraud transctions')  
plt.show()
```



Distribution of classes with time

```
# Creating fraud dataframe  
data_fraud = df[df['Class'] == 1]  
# Creating non fraud dataframe  
data_non_fraud = df[df['Class'] == 0]
```

```
# Distribution plot  
plt.figure(figsize=(8,5))  
ax = sns.distplot(data_fraud['Time'],label='fraudt',hist=False)  
ax = sns.distplot(data_non_fraud['Time'],label='non fraud',hist=False)  
ax.set(xlabel='Seconds elapsed between the transction and the first transction')  
plt.show()
```

1e-6


We do not see any specific pattern for the fraudulent and non-fraudulent transactions with respect to Time.

Hence, we can drop the Time column.

```
# Dropping the Time column
df.drop('Time', axis=1, inplace=True)
```

Distribution of classes with amount

```
# Distribution plot
plt.figure(figsize=(8,5))
ax = sns.distplot(data_fraud['Amount'],label='fraudulent',hist=False)
ax = sns.distplot(data_non_fraud['Amount'],label='non fraudulent',hist=False)
ax.set(xlabel='Transaction Amount')
plt.show()
```



We can see that the fraud transactions are mostly dense in the lower range of amount, whereas the non-fraud transactions are spreaded throughout low to high range of amount.

4. Data Preparation for Training the Model

Train-Test Split

```
# Import library
from sklearn.model_selection import train_test_split
# Putting feature variables into X
X = df.drop(['Class'], axis=1)
# Putting target variable to y
y = df['Class']
# Splitting data into train and test set 80:20
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.8, test_size=0.2, random_state=100)
```

Feature Scaling

We need to scale only the Amount column as all other columns are already scaled by the PCA transformation.

```
# Standardization method
from sklearn.preprocessing import StandardScaler
# Instantiate the Scaler
scaler = StandardScaler()

# Fit the data into scaler and transform
X_train['Amount'] = scaler.fit_transform(X_train[['Amount']])
X_train.head()
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9	
201788	2.023734	-0.429219	-0.691061	-0.201461	-0.162486	0.283718	-0.674694	0.192230	1.124319	-0.0
179369	-0.145286	0.736735	0.543226	0.892662	0.350846	0.089253	0.626708	-0.049137	-0.732566	0.0
73138	-3.015846	-1.920606	1.229574	0.721577	1.089918	-0.195727	-0.462586	0.919341	-0.612193	-0.0
208679	1.851980	-1.007445	-1.499762	-0.220770	-0.568376	-1.232633	0.248573	-0.539483	-0.813368	0.0
206534	2.237844	-0.551513	-1.426515	-0.924369	-0.401734	-1.438232	-0.119942	-0.449263	-0.717258	0.0

5 rows × 29 columns



Scale the test set

```
# Transform the test set
X_test['Amount'] = scaler.transform(X_test[['Amount']])
X_test.head()
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9	
49089	1.229452	-0.235478	-0.627166	0.419877	1.797014	4.069574	-0.896223	1.036103	0.745991	-0.1
154704	2.016893	-0.088751	-2.989257	-0.142575	2.675427	3.332289	-0.652336	0.752811	1.962566	-1.0
67247	0.535093	-1.469185	0.868279	0.385462	-1.439135	0.368118	-0.499370	0.303698	1.042073	-0.4
251657	2.128486	-0.117215	-1.513910	0.166456	0.359070	-0.540072	0.116023	-0.216140	0.680314	0.0
201903	0.558593	1.587908	-2.368767	5.124413	2.171788	-0.500419	1.059829	-0.254233	-1.959060	0.9

5 rows × 29 columns



5. Applying Models

```
# Imputing metrics
from sklearn import metrics
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classification_report, f1_score
```

```
results = pd.DataFrame(columns=['Model Name', 'Accuracy', 'F1-score', 'ROC'])
```

```
# ROC Curve function
```

```
def draw_roc( actual, probs ):
    fpr, tpr, thresholds = metrics.roc_curve( actual, probs,
```

```

drop_intermediate = False )

auc_score = metrics.roc_auc_score( actual, probs )
plt.figure(figsize=(5, 5))
plt.plot( fpr, tpr, label='ROC curve (area = %0.2f)' % auc_score )
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate or [1 - True Negative Rate]')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

return None

```

Logistic regression

```

# Importing scikit logistic regression module
from sklearn.linear_model import LogisticRegression
# Instantiate the model with best C
logistic = LogisticRegression(C=0.01)

```

```

# Fit the model on the train set
logistic_model = logistic.fit(X_train, y_train)

```

```

# Prepare results function
def display_test_results(model_name, model):

    # Prediction on the test set
    y_test_pred = model.predict(X_test)

    # Confusion matrix
    print("----- Confusion Matrix -----")
    c_matrix = metrics.confusion_matrix(y_test, y_test_pred)
    print(c_matrix)

    cm_display = ConfusionMatrixDisplay(confusion_matrix=c_matrix)
    cm_display.plot(cmap=plt.cm.Blues)
    plt.show()

    # classification_report
    print("----- classification_report -----")
    print(classification_report(y_test, y_test_pred))

    print("----- More Specific classification_report -----")
    TP = c_matrix[1,1] # true positive
    TN = c_matrix[0,0] # true negatives
    FP = c_matrix[0,1] # false positives
    FN = c_matrix[1,0] # false negatives

    # Accuracy
    print("Accuracy:-", metrics.accuracy_score(y_test, y_test_pred))

    # Sensitivity
    print("Sensitivity:-", TP / float(TP+FN))

    # Specificity
    print("Specificity:-", TN / float(TN+FP))

    # F1 score
    print("F1-Score:-", f1_score(y_test, y_test_pred))

```

```
# Predicted probability
y_test_pred_proba = model.predict_proba(X_test)[:,-1]

# roc_auc
print("----- ROC -----")
roc_auc = metrics.roc_auc_score(y_test, y_test_pred_proba)

# Plot the ROC curve
draw_roc(y_test, y_test_pred_proba)


# add all metrics score in final result store
results.loc[len(results)] = [model_name, metrics.accuracy_score(y_test, y_test_pred), f1_score(y_test, y_test_pred), roc_auc]

return None
```


Prediction results Using Logistic Regression

```
display_test_results("Logistic Regression", logistic_model)
```

```
----- Confusion Matrix -----
[[56852   14]
 [   44   52]]
```



We can see that we have very good ROC on the test set 0.98, which is almost close to 1.



XGBoost



```
# Importing XGBoost
from xgboost import XGBClassifier
params = {'learning_rate': 0.2,
          'max_depth': 2,
          'n_estimators':200,
          'subsample':0.9,
          'objective':'binary:logistic'}
```

```
# fit model on training data
xgb_model = XGBClassifier(params = params)
xgb_model.fit(X_train, y_train)
```


```
XGBClassifier(params={'learning_rate': 0.2, 'max_depth': 2, 'n_estimators': 200,
                      'objective': 'binary:logistic', 'subsample': 0.9})
```

```
Accuracy: 0.9999817772252344
```

Prediction Results Using XGBoost

```
F1-Score:- 0.6419753086419753
```

```
display_test_results("XG Boost", xgb_model)
```

Decision Tree

```
# Importing decision tree classifier
from sklearn.tree import DecisionTreeClassifier
```

```
# Model with optimal hyperparameters
decision_tree_model = DecisionTreeClassifier(criterion = "gini",
                                             random_state = 100,
                                             max_depth=5,
                                             min_samples_leaf=100,
                                             min_samples_split=100)
```

```
decision_tree_model.fit(X_train, y_train)
```


```
DecisionTreeClassifier(max_depth=5, min_samples_leaf=100, min_samples_split=100,
                      random_state=100)
```

Prediction Results Using Decision Tree

```
display_test_results("Decision Tree", decision_tree_model)
```

----- Confusion Matrix -----

```
[[56836  30]
 [  40   56]]
```



----- classification_report -----

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56866
1	0.65	0.58	0.62	96
accuracy			1.00	56962
macro avg	0.83	0.79	0.81	56962
weighted avg	1.00	1.00	1.00	56962

Random Forest

Sensitivity:- 0.5833333333333334

```
# Importing random forest classifier
from sklearn.ensemble import RandomForestClassifier
random_forest_model = RandomForestClassifier(bootstrap=True,
                                             max_depth=5,
                                             min_samples_leaf=50,
                                             min_samples_split=50,
                                             max_features=10,
                                             n_estimators=100)
```

```
# Fit the model
random_forest_model.fit(X_train, y_train)
```


```
RandomForestClassifier(max_depth=5, max_features=10, min_samples_leaf=50,
                       min_samples_split=50)
```

Prediction Results Using Random forest

```
display_test_results("Random Forest", random_forest_model)
```

----- Confusion Matrix -----

```
[[56844  22]
 [  33  63]]
```



----- classification_report -----

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56866
1	0.74	0.66	0.70	96
accuracy			1.00	56962
macro avg	0.87	0.83	0.85	56962
weighted avg	1.00	1.00	1.00	56962

----- More Specific classification_report -----

Accuracy:- 0.9990344440153085
 Sensitivity:- 0.65625
 Specificity:- 0.9996131255935005
 F1-Score:- 0.6961325966850829

----- ROC -----

Receiver operating characteristic example




6. Summary

0.6

```
results.sort_values(by="ROC", ascending=False)
```

	Model Name	Accuracy	F1-score	ROC	
0	Logistic Regression	0.998982	0.641975	0.977696	
1	XG Boost	0.999350	0.797814	0.976238	
3	Random Forest	0.999034	0.696133	0.961433	
2	Decision Tree	0.998771	0.615385	0.921750	

We can see that **XG Boost** Algorithm performing Better.

 0s completed at 2:00 PM

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.